


180040165-DALALI ARIF
180040668-V.V.N.A.VISHAL
180040670-ANUSH NAIDHU

Guided by: GOPI TIAK VESALA

PROJECT TITLE

- ▼ Iris data set: Predict the chess of the flower based on available attributes.



Iris Versicolor



Iris Setosa



Iris Virginica

- iris versicolor** is a flowering herbaceous perennial plant, growing 10–80 cm (4–31 in) high. It tends
- ▼ to form large clumps from thick, creeping rhizomes. The unwinged, erect stems generally have basal leaves that are more than 1 cm ($1/2$ in) wide.

Iris setosa, the bristle-pointed iris, is a species of flowering plant in the genus *Iris* of the family Iridaceae, it belongs to the subgenus *Limniris* and the series *Tripetalae*.

Iris virginica, with the common name Virginia iris, is a perennial species of flowering plant, native to eastern North America. It is common

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from pandas.plotting import parallel_coordinates
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn import metrics
from sklearn.naive_bayes import GaussianNB
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis, QuadraticDiscriminantAnalysis
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
```

```
data = pd.read_csv('/Iris (1).csv')
```

```
data.head(5)
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa

```
data.describe()
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	75.500000	5.843333	3.054000	3.758667	1.198667
std	43.445368	0.828066	0.433594	1.764420	0.763161
min	1.000000	4.300000	2.000000	1.000000	0.100000
25%	38.250000	5.100000	2.800000	1.600000	0.300000
50%	75.500000	5.800000	3.000000	4.350000	1.300000
75%	112.750000	6.400000	3.300000	5.100000	1.800000
max	150.000000	7.900000	4.400000	6.900000	2.500000

```
data.groupby('Species').size()
```

```
Species
Iris-setosa      50
Iris-versicolor  50
Iris-virginica   50
dtype: int64
```

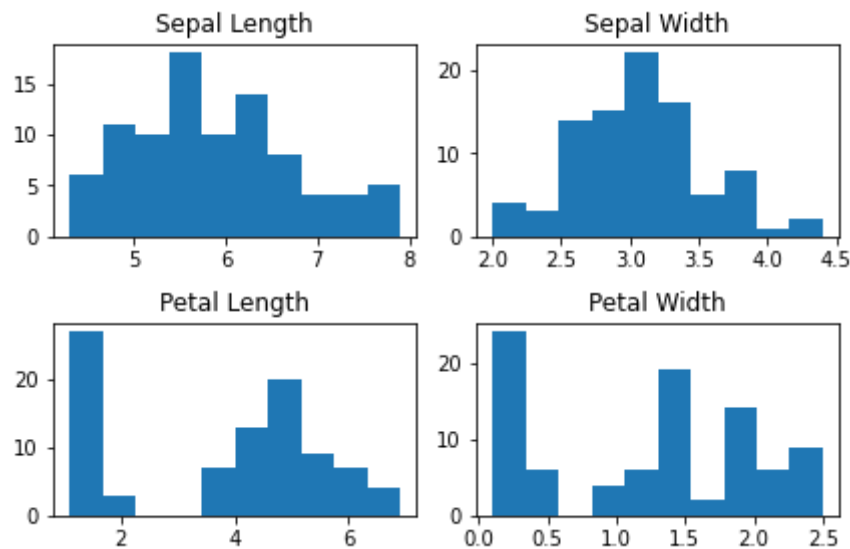
```
train, test = train_test_split(data, test_size = 0.4, stratify = data['Species'], random_state = 42)
```

```
n_bins = 10
fig, axs = plt.subplots(2, 2)
axs[0,0].hist(train['SepalLengthCm'], bins = n_bins);
axs[0,0].set title('Sepal Length');
```

```

fig, axs = plt.subplots(2, 2)
axs[0,1].hist(train['SepalWidthCm'], bins = n_bins);
axs[0,1].set_title('Sepal Width');
axs[1,0].hist(train['PetalLengthCm'], bins = n_bins);
axs[1,0].set_title('Petal Length');
axs[1,1].hist(train['PetalWidthCm'], bins = n_bins);
axs[1,1].set_title('Petal Width');
# add some spacing between subplots
fig.tight_layout(pad=1.0);

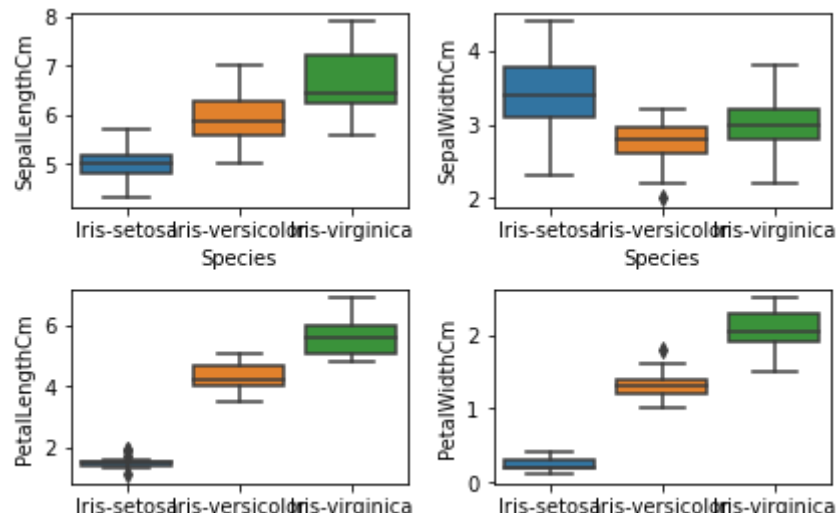
```




```

fig, axs = plt.subplots(2, 2)
fn = ["SepalLengthCm", "SepalWidthCm", "PetalLengthCm", "PetalWidthCm"]
cn = ['Iris-setosa', 'Iris-versicolor', 'Iris-virginica']
sns.boxplot(x = 'Species', y = 'SepalLengthCm', data = train, order = cn, ax = axs[0,0]);
sns.boxplot(x = 'Species', y = 'SepalWidthCm', data = train, order = cn, ax = axs[0,1]);
sns.boxplot(x = 'Species', y = 'PetalLengthCm', data = train, order = cn, ax = axs[1,0]);
sns.boxplot(x = 'Species', y = 'PetalWidthCm', data = train, order = cn, ax = axs[1,1]);
# add some spacing between subplots
fig.tight_layout(pad=1.0);

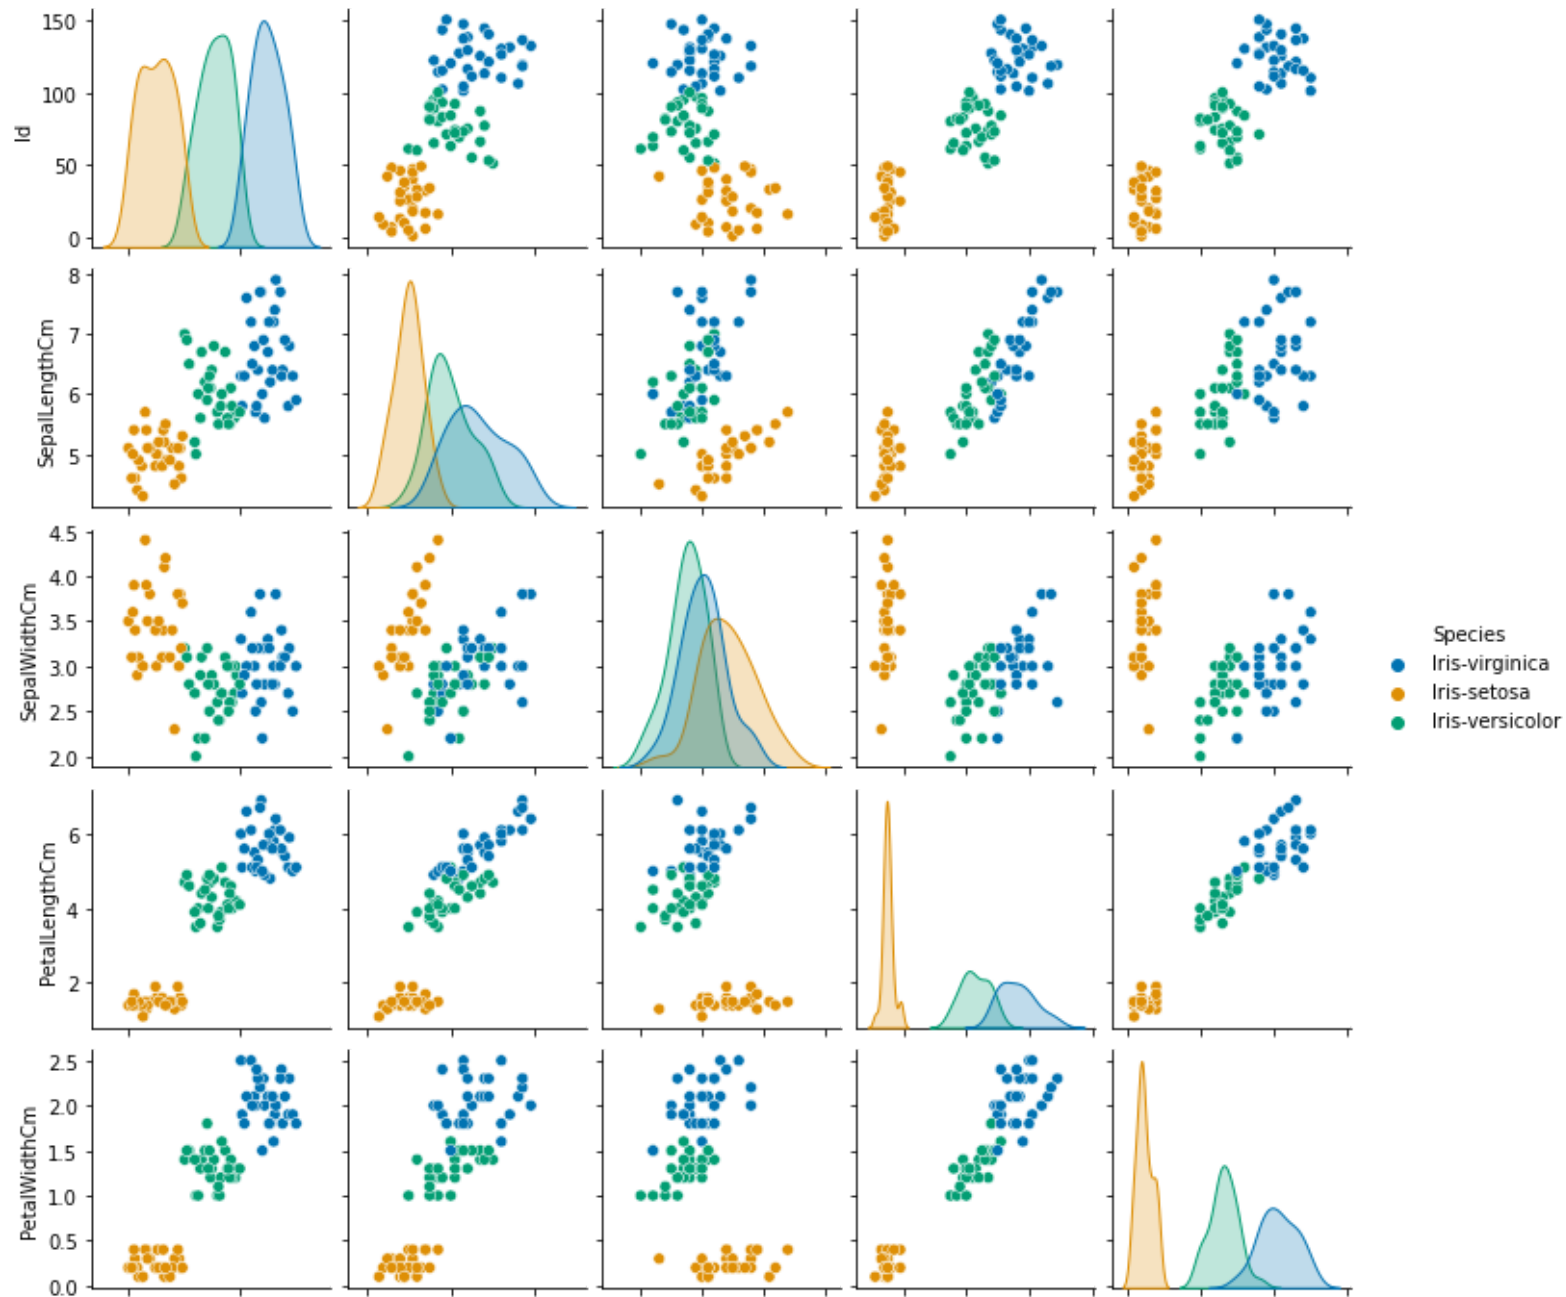
```




```
sns.violinplot(x="Species", y="PetalLengthCm", data=train, size=5, order = cn, palette = 'colorblind');
```



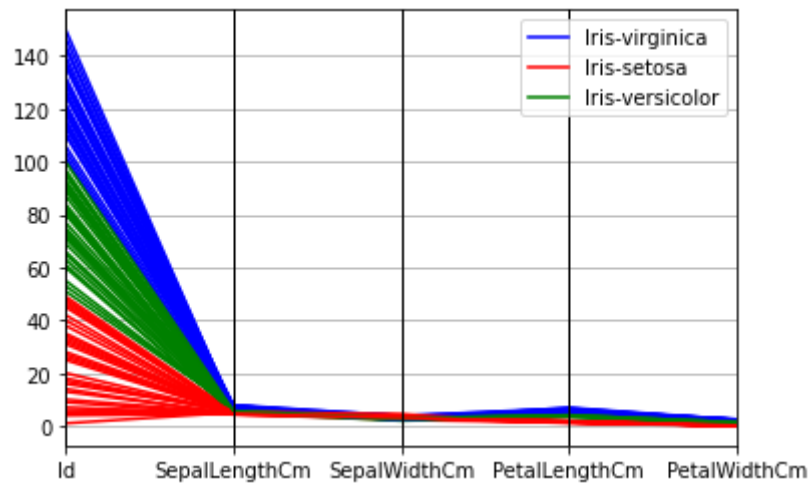
```
sns.pairplot(train, hue="Species", height = 2, palette = 'colorblind');
```



```
corrmat = train.corr()
sns.heatmap(corrmat, annot = True, square = True);
```



```
parallel_coordinates(train, "Species", color = ['blue', 'red', 'green']);
```



```
X_train = train[['SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm']]
y_train = train.Species
```

```
X_test = test[['SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm']]
y_test = test.Species

mod_dt = DecisionTreeClassifier(max_depth = 3, random_state = 1)
mod_dt.fit(X_train,y_train)
prediction=mod_dt.predict(X_test)
print('The accuracy of the Decision Tree is', "{:.3f}".format(metrics.accuracy_score(prediction,y_test)))
```

The accuracy of the Decision Tree is 0.983

```
mod_dt.feature_importances_

array([0.          , 0.          , 0.42430866, 0.57569134])
```

```
plt.figure(figsize = (10,8))
plot_tree(mod_dt, feature_names = fn, class_names = cn, filled = True);
```



```


PetalWidthCm <= 0.7
  gini = 0.667
  samples = 90
  value = [30, 30, 30]
  class = Iris-setosa

```

```

disp = metrics.plot_confusion_matrix(mod_dt, X_test, y_test,
                                     display_labels=cn,
                                     cmap=plt.cm.Blues,
                                     normalize=None)
disp.ax_.set_title('Decision Tree Confusion matrix, without normalization');

```



```

from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report

```

```

X = data.iloc[:, :-1].values
y = data.iloc[:, -1].values

```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0)

```

```

from sklearn.linear_model import LogisticRegression
classifier = LogisticRegression()
classifier.fit(X_train, y_train)
y_pred = classifier.predict(X_test)
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
from sklearn.metrics import accuracy_score
print('accuracy is', accuracy_score(y_pred, y_test))

```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	11
Iris-versicolor	1.00	1.00	1.00	13
Iris-virginica	1.00	1.00	1.00	6
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```

[[11  0  0]
 [ 0 13  0]
 [ 0  0  6]]

```

accuracy is 1.0

/usr/local/lib/python3.6/dist-packages/sklearn/linear_model/_logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1): STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

Naive Bayes

```

from sklearn.naive_bayes import GaussianNB
classifier = GaussianNB()
classifier.fit(X_train, y_train)

```

```

y_pred = classifier.predict(X_test)

# Summary of the predictions made by the classifier
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
# Accuracy score
from sklearn.metrics import accuracy_score
print('accuracy is',accuracy_score(y_pred,y_test))

```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	11
Iris-versicolor	1.00	1.00	1.00	13
Iris-virginica	1.00	1.00	1.00	6
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```

[[11  0  0]
 [ 0 13  0]
 [ 0  0  6]]
accuracy is 1.0

```

```

from sklearn.svm import SVC

```

```

classifier = SVC()
classifier.fit(X_train, y_train)

```

```

y_pred = classifier.predict(X_test)

```

```

# Summary of the predictions made by the classifier
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
# Accuracy score
from sklearn.metrics import accuracy_score
print('accuracy is',accuracy_score(y_pred,y_test))

```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	11
Iris-versicolor	1.00	1.00	1.00	13
Iris-virginica	1.00	1.00	1.00	6
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
[[11  0  0]
 [ 0 13  0]
 [ 0  0  6]]
accuracy is 1.0
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
classifier = KNeighborsClassifier(n_neighbors=8)
classifier.fit(X_train, y_train)
```

```
y_pred = classifier.predict(X_test)
```

```
# Summary of the predictions made by the classifier
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
# Accuracy score
from sklearn.metrics import accuracy_score
print('accuracy is', accuracy_score(y_pred, y_test))
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	11
Iris-versicolor	1.00	1.00	1.00	13
Iris-virginica	1.00	1.00	1.00	6
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
[[11  0  0]
 [ 0 13  0]
 [ 0  0  6]]
accuracy is 1.0
```

```
dict={'S.No':[1,2,3,4,5],'Model':['Decision Tree','LogisticRegression','GaussianNB','Support Vector Machine','K Nearest Neighbours'].
```

```
dict
```

```
{'Accuracy': [0.983, 1.0, 1.0, 1.0, 1.0],
 'Model': ['Decision Tree',
 'LogisticRegression',
 'GaussianNB',
 'Support Vector Machine',
 'K Nearest Neighbours'],
 'S.No': [1, 2, 3, 4, 5]}
```

```
data=pd.DataFrame.from_dict(dict)
```

```
data
```

	S.No	Model	Accuracy
0	1	Decision Tree	0.983
1	2	LogisticRegression	1.000
2	3	GaussianNB	1.000
3	4	Support Vector Machine	1.000
4	5	K Nearest Neighbours	1.000

